

# DATA STRUCTURES & ALGORITHMS

## Lecture 5: Advanced Data Structures B-Trees

---

Lecturer: Dr. Nguyen Hai Minh  
[nhminh@fit.hcmus.edu.vn](mailto:nhminh@fit.hcmus.edu.vn)

# OUTLINE

- ❑ Introduction
- ❑ Basic operations on B-Tree
- ❑ Deleting a key from a B-Tree

# Motivatio of B-Trees

- ❑ So far, we have assumed that we can store an entire data structure in main memory
- ❑ What if we have so much data that it won't fit?
- ❑ Storing it on disk requires different approach to efficiency
- ❑ Assuming that a disk spins at 3600 RPM, one revolution occurs in  $1/60$  of a second, or 16.7ms
- ❑ Crudely speaking, one disk access takes about the same time as 200,000 instructions



# Motivation of B-Trees

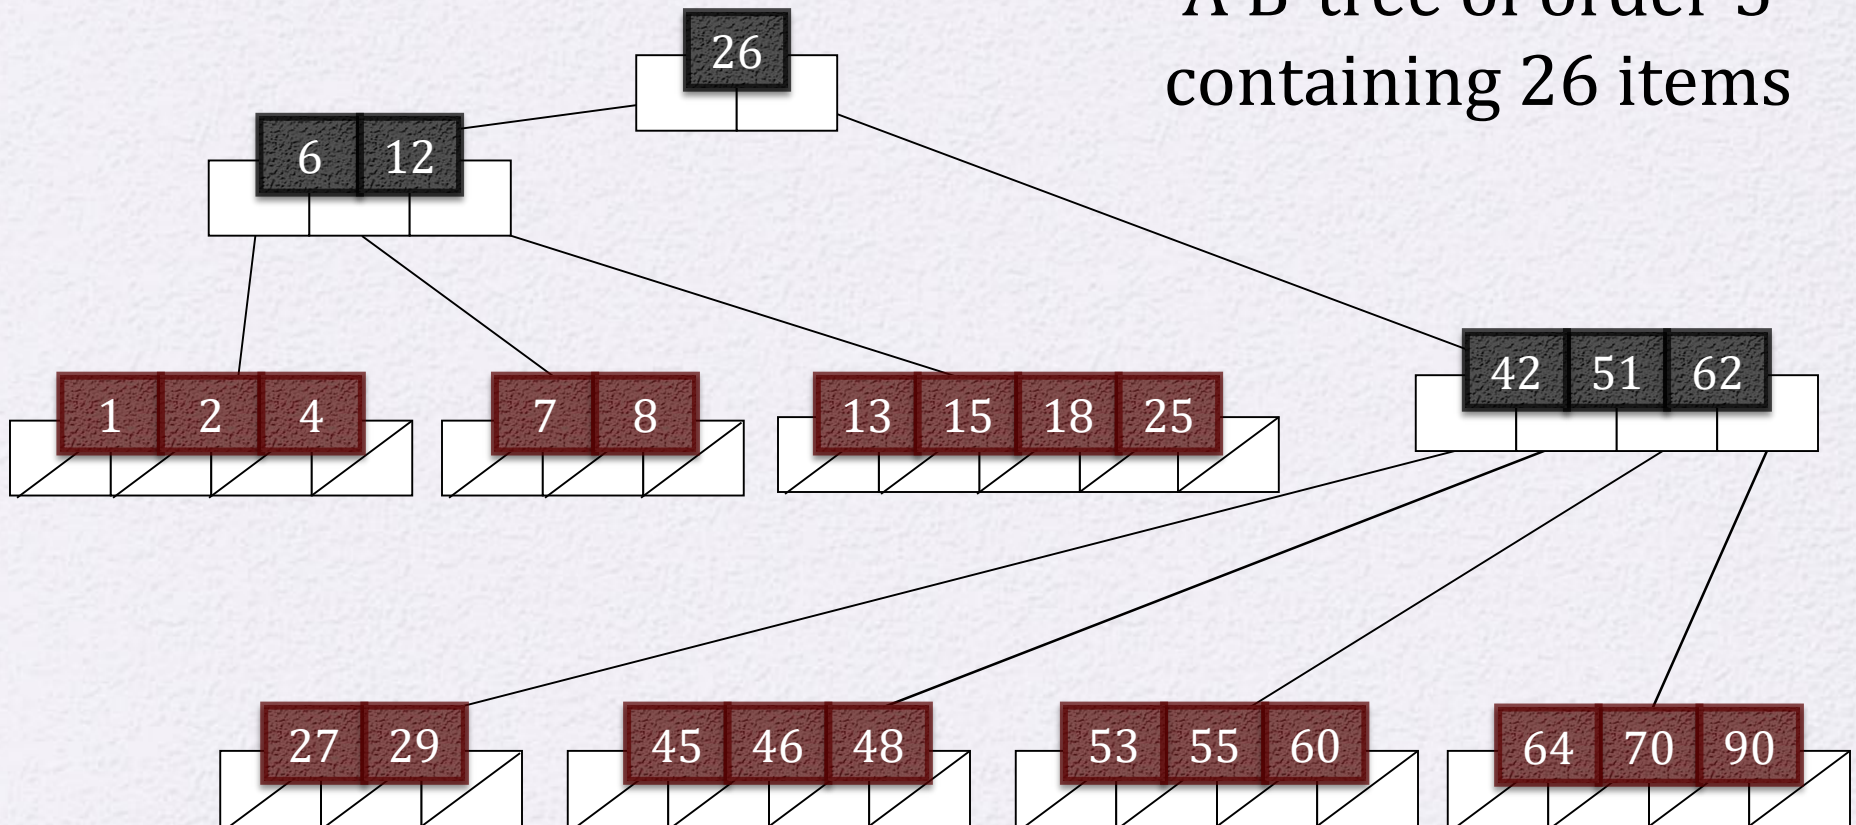
- ❑ Assume that we use an AVL tree to store about 20 million records
- ❑ We end up with a very deep binary tree with lots of different disk accesses;  $\log_2 20,000,000 \approx 24$ , so this takes about **0.2** seconds
- ❑ We know we can't improve on the  $\log_2 n$  lower bound on search for a binary tree
- ❑ But, the solution is to use more branches and thus reduce the height of the tree!
- ❑ As branching increases, depth decreases

# Definition of a B-Tree

- ❑ A B-tree of order  $m$  is an  $m$ -way tree (i.e., a tree where each node may have up to  $m$  children) in which:
  1. The number of keys in each non-leaf node is one less than the number of its children and these keys partition the keys in the children in the fashion of a search tree
  2. All leaves are on the same level
  3. All non-leaf nodes except the root have at least  $\lceil m/2 \rceil$  children
  4. The root is either a leaf node, or it has from 2 to  $m$  children
  5. A leaf node contains  $\lceil m/2 \rceil - 1$  to  $m - 1$  keys
- ❑ The number  $m$  should always be odd

# An example B-Tree

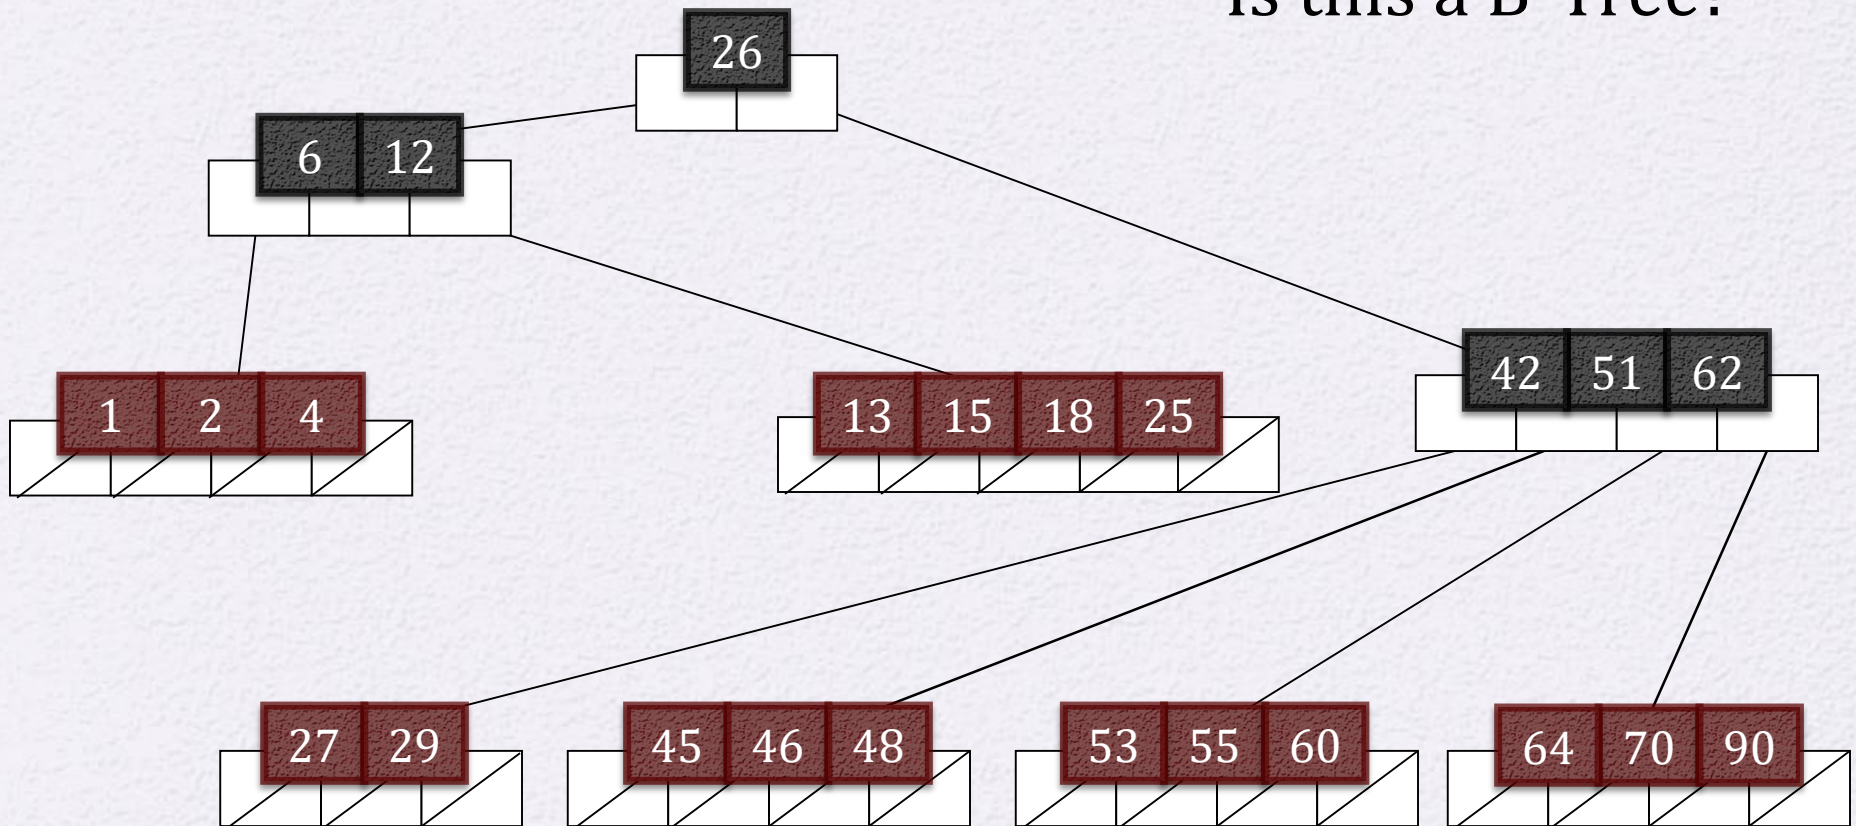
A B-tree of order 5  
containing 26 items





# An example B-Tree

Is this a B-Tree?



# B-Tree's maximum height

□ Let  $h$  be the height of B-Tree

→ The maximum number of nodes in a  $m$ -order B-Tree is:

$$m^{h+1} - 1$$



# B-Tree's maximum height

- ❑ The maximum number of items in a B-tree of order  $m$  and height  $h$ :

root	$m - 1$
level 1	$m(m - 1)$
level 2	$m^2(m - 1)$
...	
level $h$	$m^h(m - 1)$

- ❑ So, the total number of items is

$$(1 + m + m^2 + m^3 + \dots + m^h)(m - 1) =$$
$$[(m^{h+1} - 1) / (m - 1)] (m - 1) = \mathbf{m^{h+1} - 1}$$

- ❑ When  $m = 5$  and  $h = 2$  this gives  $5^3 - 1 = 124$

# Inserting into a B-Tree

- ❑ Attempt to insert the new key into a leaf
- ❑ If **leaf** becomes too big,
  - Split the leaf into two
  - Promoting the middle key to the leaf's parent
- ❑ If the **parent** becomes too big
  - Split the parent into two
  - Promoting the middle key
- ❑ This strategy might have to be repeated all the way to the top
- ❑ If necessary, the root is split in two and the middle key is promoted to a new root, *making the tree one level higher*

# Inserting into a B-Tree

## Example

- ❑ Suppose we start with an empty B-tree and keys arrive in the following order: 1 12 8 2 25 5 14 28 17 7 52 16 48 68 3 26 29 53 55 45
- ❑ We want to construct a B-tree of order 5



# Inserting into a B-Tree

## Example

1 12 8 2 25 5 14  
28 17 7 52 16 48  
68 3 26 29 53 55  
45





Insert 1      Insert 12      Insert 8      Insert 2

# Inserting into a B-Tree

## Example

1 12 8 2 **25** 5 14  
28 17 7 52 16 48  
68 3 26 29 53 55  
45




Insert **25**

# Inserting into a B-Tree

## Example

1 12 8 2 25 5 14  
28 17 7 52 16 48  
68 3 26 29 53 55  
45




Insert 5, 14, 28



# Inserting into a B-Tree

## Example



1 12 8 2 25 5 14  
28 17 7 52 16 48  
68 3 26 29 53 55  
45



# Inserting into a B-Tree

## Example

1 12 8 2 25 5 14  
28 17 7 52 16 48  
68 3 26 29 53 55  
45





Insert 7, 52, 16, 48

# Inserting into a B-Tree

## Example

1 12 8 2 25 5 14  
28 17 7 52 16 48  
**68** 3 26 29 53 55  
45





Insert 68



# Inserting into a B-Tree

## Example

1 12 8 2 25 5 14  
28 17 7 52 16 48  
68 3 26 29 53 55  
45




Insert 3

# Inserting into a B-Tree

## Example

1 12 8 2 25 5 14  
28 17 7 52 16 48  
68 3 26 29 53 55  
45

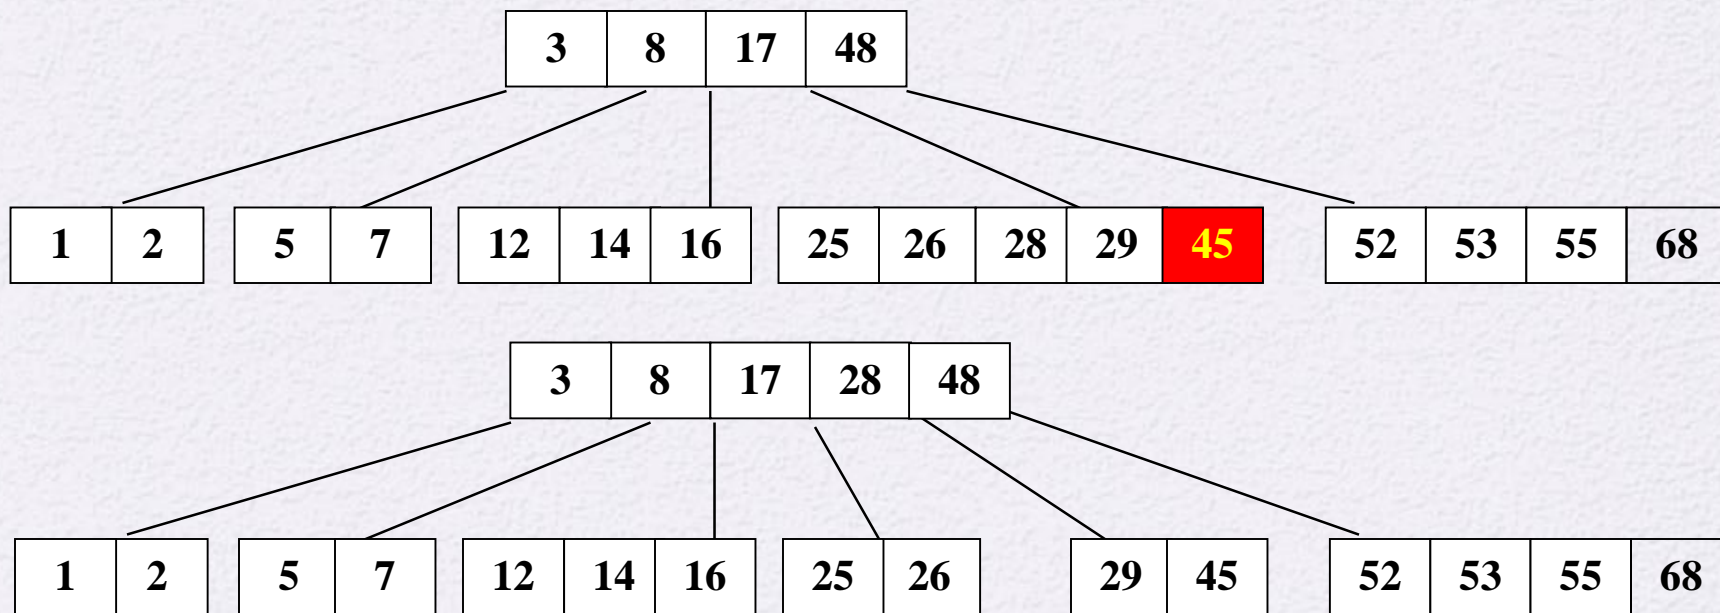


Insert **26, 29, 53, 55**

# Inserting into a B-Tree

## Example

1 12 8 2 25 5 14  
28 17 7 52 16 48  
68 3 26 29 53 55  
45





Insert 45



# Inserting into a B-Tree

## Example

1 12 8 2 25 5 14  
28 17 7 52 16 48  
68 3 26 29 53 55  
45



Insert 45

# Removal from a B-tree

❑ Remove a key  $k$  from a leaf.

1. If  $k$  is in a leaf node, and removing it doesn't cause that leaf node to have too few keys, then simply remove  $k$ .
2. If  $k$  is NOT in a leaf then it is guaranteed (by the nature of a B-tree) that its predecessor or successor will be in a leaf -- in this case we can delete  $k$  and promote the predecessor or successor of  $k$  to  $k$ 's position.


# Removal from a B-tree (2)

- ❑ (1) & (2) may lead to a leaf node  $L$  has less than min. number of keys
- ➔ Look at the siblings immediately adjacent to the leaf
  - 3: if one of them has more than the min. number of keys then we can promote one of its keys to the parent and take the parent key into  $L$
  - 4: otherwise, combine  $L$  and one of its neighbours with their shared parent, repeat the process up to the root, if required



# Type #1: Simple leaf deletion


Assuming a 5-way  
B-Tree, as before...



Delete 2: Since there are enough  
keys in the node, just delete it


*Note when printed: this slide is animated*

# Type #2: Simple non-leaf deletion



*Note when printed: this slide is animated*


# Type #4: Too few keys in node and its siblings



*Note when printed: this slide is animated*




# Type #4: Too few keys in node and its siblings




*Note when printed: this slide is animated*

# Type #3: Enough siblings



*Note when printed: this slide is animated*

# Type #3: Enough siblings



*Note when printed: this slide is animated*



# Exercise in Removal from a B-Tree

❑ Given 5-way B-tree created by these data:

3, 7, 9, 23, 45, 1, 5, 14, 25, 24, 13, 11, 8, 19, 4, 31, 35,  
56

❑ Add these further keys: 2, 6, 12

❑ Delete these keys: 4, 5, 7, 3, 14

# Reasons for using B-Trees

- ❑ When searching tables held on disc, the cost of each disc transfer is high but doesn't depend much on the amount of data transferred, especially if consecutive items are transferred
  - If we use a B-tree of order 101, say, we can transfer each node in one disc read operation
  - A B-tree of order 101 and height 3 can hold  $101^4 - 1$  items (approximately 100 million) and any item can be accessed with 3 disc reads (assuming we hold the root in memory)
- ❑ If we take  $m = 3$ , we get a **2-3 tree**, in which non-leaf nodes have two or three children (i.e., one or two keys)
  - B-Trees are always balanced (since the leaves are all at the same level), so 2-3 trees make a good type of balanced tree

# Comparing Trees

## ❑ Binary trees

- Can become *unbalanced* and *lose* their good time complexity (big O)
- AVL trees are strict binary trees that *overcome the balance problem*
- Heaps remain balanced but only *prioritise* (not order) the keys

## ❑ Multi-way trees

- B-Trees can be *m*-way, they can have any (odd) number of children
- One B-Tree, the 2-3 (or 3-way) B-Tree, *approximates* a permanently balanced binary tree, exchanging the AVL tree's balancing operations for insertion and (more complex) deletion operations



# Q&A